

SQL Server 2012 – Database Development

Lesson 1: Database
Objects: Indexes &
Views



Lesson Objectives

➤ In this lesson, you will learn:

- Creating Indexes
- Querying the sysindexes Table
- Performance Considerations
- Creating Views
- Creating Sequence





Index – An Overview

- Database systems generally use indexes to provide fast access to relational data
- An index is a separate physical data structure that enables queries to access one or more data rows fast
- This structure is known as B-Tree Structure
- Proper tuning of index is therefore a key for query performance
- Database Engine uses index to find the data just like one uses index in a book
- When a table is dropped , indexes also get dropped automatically
- Only the owner of the table can create indexes
- SQL Server supports two types of indexes
 - Clustered
 - Non clustered



How SQL Server access data?

- SQL Server accesses data in one of two ways:
- By scanning all the data pages in a table, which is called a table scan. When SQL Server performs a table scan, it:
 - Starts at the beginning of the table
 - Scans from page to page through all the rows in the table
 - Extracts the rows that meet the criteria of the query
- By using indexes. When SQL Server uses an index, it:
 - Traverses the index tree structure to find rows that the query requests
 - Extracts only the needed rows that meet the criteria of the query



Clustered Index

- A clustered index determines the physical order of the data in a table
- Database Engine allows the creation of a single clustered index per table
- If a clustered index is defined for a table, the table is called a clustered table
- A Unique Clustered index is built by default for each table, for which you define the primary key using the primary key constraint
- Also, each clustered index is unique by default that is, each data value can appear only once in a column for which the clustered index is defined



Non-Clustered Index

- A Non-Clustered index has the same index structure as a clustered index
- A Non-Clustered index does not change the physical order of the rows in the table
- A table can have more than one non clustered index
- Unique Non-Clustered index will be created automatically when you create unique key on a column to enforce uniqueness of key value



Filtered Index

- SQL Server 2008 introduces filtered indexes and statistics
- The Non-Clustered indexes now can be created based on a predicate, and only the subset of rows for which the predicate holds true are stored in the index B-Tree
- Well-designed filtered indexes can improve query performance and plan quality because they are smaller than nonfiltered indexes
- We can also reduce index maintenance cost by using filtered indexes because there is less data to maintain
- Filtered indexes also obviously reduce storage costs

```
USE AdventureWorks
GO
CREATE NONCLUSTERED INDEX idx_currate_notnull
ON Sales.SalesOrderHeader(CurrencyRateID)
WHERE CurrencyRateID IS NOT NULL
```



Creating and Dropping Indexes

- Indexes are created automatically on tables with PRIMARY KEY or UNIQUE constraints
- Indexes can also be created using the CREATE INDEX Statement

```
USE Northwind  
CREATE CLUSTERED INDEX CL_lastname  
ON employees(lastname)
```

- Indexes can be dropped using the DROP command

```
USE Northwind  
DROP INDEX employees.CL_lastname
```




Creating and Dropping Indexes

- To create non clustered index ncl_deptno

```
USE Northwind
CREATE NON CLUSTERED INDEX NCL_deptno
ON employees(deptno)
```

- Using the DROP INDEX Statement

```
USE Northwind
DROP INDEX employees.NCL_deptno
```



Creating Unique Indexes

- Unique index can be non clustered or clustered
- Unique non clustered index is automatically created when a column has UNIQUE constraint
- Unique Clustered index is automatically created when column has a PRIMARY KEY constraint
- Ensures column(s) have unique value
- There is no difference in the way Unique constraint and Unique index work , except for syntax

```
USE Northwind
CREATE UNIQUE NONCLUSTERED INDEX U_CustID
ON customers(CustomerID)
```



Creating Composite Indexes

- Index of two /more columns are said to be composite

USE Northwind

```
CREATE UNIQUE NONCLUSTERED INDEX U_OrdID_ProdID  
ON [Order Details] (OrderID, ProductID)
```

<i>Order Details</i>				
<i>OrderID</i>	<i>ProductID</i>	<i>UnitPrice</i>	<i>Quantity</i>	<i>Discount</i>
10248	11	14.000	12	0.0
10248	42	9.800	10	0.0
10248	72	34.800	5	0.0

Column 1

Column 2

Composite Key



Columnstore Indexes

- SQL Server 2012 introduces ColumnStore Indexes.

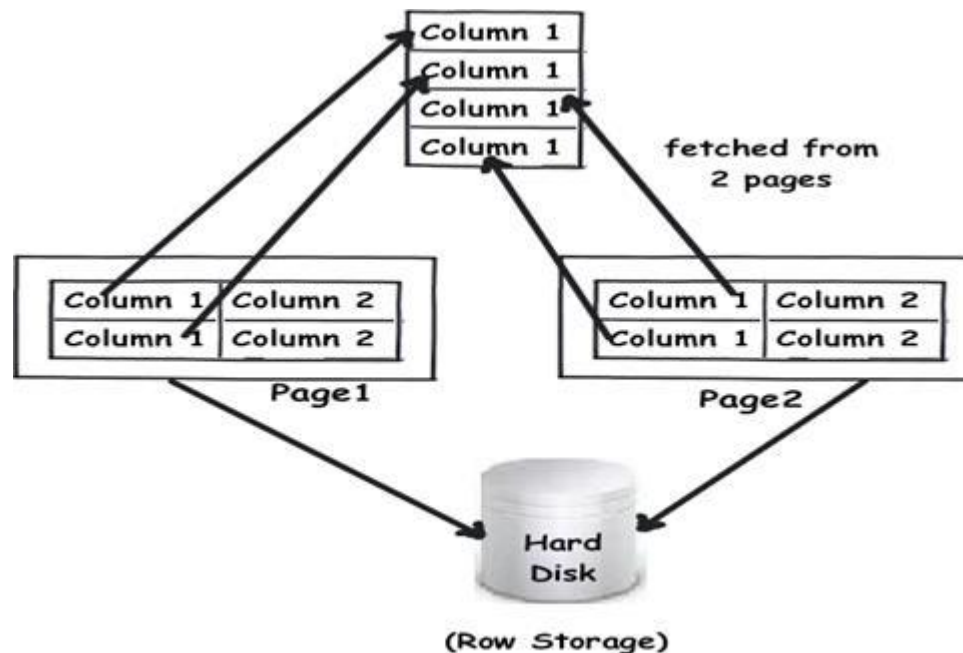
Benefits of using SQL Server ColumnStore Indexes

- Faster query performance for common data warehouse queries as only required columns/pages in the query are fetched from disk
- Data is stored in a highly compressed form to reduce the storage space
- Frequently accessed columns (pages that contains data for these columns) remain in memory because a high ratio of compression is used in the pages and less pages are involved

Columnstore Indexes



- Relational database store data "row wise". These rows are further stored in 8 KB page size.
- For instance you can see in the below figure we have table with two columns "Column1" and "Column2". You can see how the data is stored in two pages i.e. "page1" and "page2". "Page1" has two rows and "page2" also has two rows.






Columnstore Indexes

- If you want to fetch only "column1", you have to pull records from two pages i.e. "Page1" and "Page2".
- As we have to fetch data from two pages its bit performance intensive.

Columnstore Indexes



- If somehow we can store data column wise we can avoid fetching data from multiple pages.
- That's what column store indexes do.
- When you create a column store index it stores same column data in the same page. You can see from the diagram (shown below), we now need to fetch "column1" data only from one page rather than querying multiple pages.



Limitations of SQL Server ColumnStore Indexes



There are several limitations of using SQL Server ColumnStore indexes over Row Store indexes including:

- A table with a ColumnStore Index cannot be updated
- ColumnStore index creation takes more time (1.5 times almost) than creating a B-tree index (on same set of columns) because the data is compressed
- A table can have only one ColumnStore Index and hence you should consider including all columns or at least all those frequently used columns of the table in the index
- A ColumnStore Index can only be non cluster and non unique index; you cannot specify ASC/DESC or INCLUDE clauses
- Not all data types (binary, varbinary, image, text, ntext, varchar(max), nvarchar(max), etc.) are supported
-



Limitations of SQL Server ColumnStore Indexes

- The definition of a ColumnStore Index cannot be changed with the ALTER INDEX command, you need to drop and create the index or disable it then rebuild it
- You can create a ColumnStore index on a table which has compression enabled, but you cannot specify the compression setting for the column store index
- A ColumnStore Index cannot be created on view
- A ColumnStore Index cannot be created on table which uses features like Replication, Change Tracking, Change Data Capture and Filestream



Columnstore Indexes

➤ CREATE NONCLUSTERED COLUMNSTORE INDEX idx_colSale
ON myTable (OrderDate, ProductID, SaleAmount)



Obtaining information on Indexes

- Using the `sp_helpindex` System Stored Procedure

```
USE Northwind  
EXEC sp_helpindex Customers
```

- Using the `sp_help tablename` System Stored Procedure



Indexes – Performance Considerations

- Create indexes on foreign keys
- Create the clustered index before nonclustered indexes
- Consider before creating composite indexes
- Create multiple indexes for a table that is read frequently
- Use the index tuning wizard get statics of index usage

Demo



➤ Creating Indexes






Views – An Overview

- Views are Virtual tables, which provides access to a subset of columns from one or more tables
- Created from one or more base tables or other views
- Internally Views are stored queries
- Views are created when
 - To hide the complexity of the underlying database schema, or customize the data and schema for a set of users.
 - To control access to rows and columns of data.
- Objective of creating views is Abstraction , not performance



Views – An Overview

<i>Employees</i>			
<i>EmployeeID</i>	<i>LastName</i>	<i>Firstname</i>	<i>Title</i>
1	Davolio	Nancy	~~~
2	Fuller	Andrew	~~~
3	Leverling	Janet	~~~



```
USE Northwind
GO
CREATE VIEW dbo.EmployeeView
AS
SELECT LastName, Firstname
FROM Employees
```

<i>EmployeeView</i>	
<i>Lastname</i>	<i>Firstname</i>
Davolio	Nancy
Fuller	Andrew
Leverling	Janet





Views – Advantages

- Focus the Data for Users
 - Focus on important or appropriate data only
 - Limit access to sensitive data
- Mask Database Complexity
 - Hide complex database design
 - Simplify complex queries, including distributed queries to heterogeneous data
- Simplify Management of User Access on Data



Views – Types

- Standard Views
- Indexed Views
- Partitioned Views




Defining Views

- Creating views
- Altering and dropping views
- Locating view definition information
- Hiding view definitions



Creating Views

➤ Creating a View



```
CREATE VIEW dbo.OrderSubtotalsView (OrderID, Subtotal)
AS
SELECT OD.OrderID,
       SUM(CONVERT(money,(OD.UnitPrice*Quantity*(1-Discout)/100))*100)
FROM [Order Details] OD
GROUP BY OD.OrderID
GO
```

➤ Restrictions on View Definitions

- Cannot include ORDER BY clause
- Cannot include INTO keyword



Example – Views with Join Query

Orders

OrderID	CustomerID	RequiredDate	ShippedDate
10663	BONAP	1997-09-24	1997-10-03
10827	BONAP	1998-01-26	1998-02-06
10427	PICCO	1997-02-24	1997-03-03
10451	QUICK	1997-03-05	1997-03-12
10515	QUICK	1997-05-07	1997-05-23

Customers

CustomerID	CompanyName	ContactName
BONAP	Bon app'	Laurence Lebihan
PICCO	Piccolo und mehr	Georg Pippis
QUICK	QUICK-Stop	Horst Kloss

```
USE Northwind
GO
CREATE VIEW dbo.ShipStatusView
AS
SELECT OrderID, RequiredDate,
ShippedDate, ContactName
FROM Customers
INNER JOIN Orders
ON Customers.CustomerID = Orders.CustomerID
WHERE RequiredDate < ShippedDate
```

ShipStatusView

OrderID	ShippedDate	ContactName
10264	1996-08-23	Laurence Lebihan
10271	1996-08-30	Georg Pippis
10280	1996-09-12	Horst Kloss



Altering & Dropping Views

➤ Altering Views

- Retains assigned permissions
- Causes new SELECT statement and options to replace existing definition

```
USE Northwind
GO
ALTER VIEW dbo.EmployeeView
AS
SELECT LastName, FirstName, Extension
FROM Employees
```

➤ Dropping Views

```
DROP VIEW dbo.ShipStatusView
```



Locating View Definition Information

➤ Locating View Definitions

- Not available if view was created using WITH ENCRYPTION option

➤ Locating View Dependencies

- Lists objects upon which view depends
- Lists objects that depend on a view



Hiding View Definition

- Use the WITH ENCRYPTION Option
- Do not delete entries in the syscomments table

```
USE Northwind
GO
CREATE VIEW dbo.[Order Subtotals]
    WITH ENCRYPTION
AS
SELECT OrderID,
    Sum(CONVERT(money,(UnitPrice*Quantity*(1-Discount)/100))*100)
    AS Subtotal
FROM [Order Details]
GROUP BY OrderID
GO
```



Modifying Data through View

- Cannot affect more than one underlying table
- Cannot be made to columns having aggregation
- Depends on the constraints placed on the base tables
- Are verified if the WITH CHECK OPTION has been specified



Views – Recommended Practices

- Use a Standard Naming Convention
- dbo Should Own All Views
- Verify Object Dependencies Before You Drop Objects
- Never Delete Entries in the syscomments Table
- Carefully Evaluate Creating Views Based on Views

Demo



➤ Working with Views





Summary

➤ In this lesson, you have learnt:

- Creating Indexes
- Types of indexes
 - Clustered Index, Non clustered index ,Filtered Indexes ,Column store Indexes
- Creating and modifying Views





Review Question

- Question 1: ----- Gets created automatically for Primary key constrain
 - clustered index
 - Unique clustered index
 - Unique Non clustered index

- Question 2: ----- option with views will not stored base query of views in syscomments table

- Question 3: A table can have multiple unique non clustered index
 - True/False

